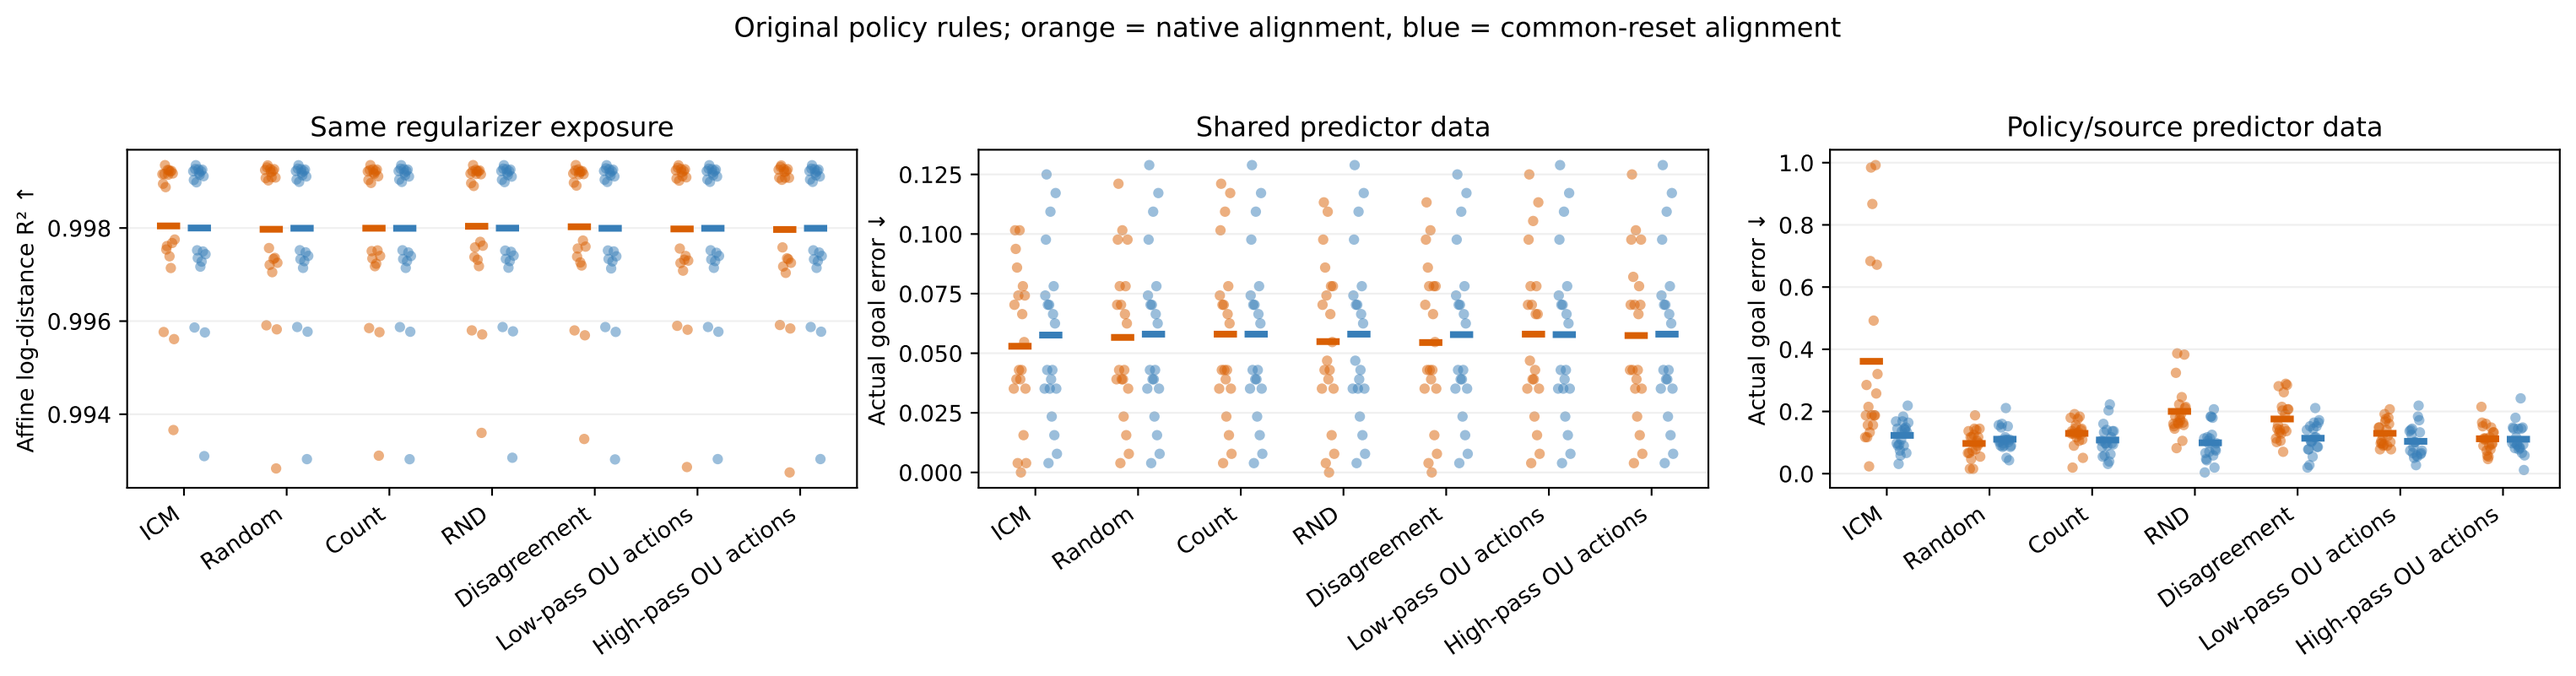
Original policy rules; orange = native alignment, blue = common-reset alignment
Same regularizer exposure
Shared predictor data
Policy/source predictor data
Affine log-distance R² ↑
Actual goal error ↓
Actual goal error ↓
0.998
0.996
0.994
0.125
0.100
0.075
0.050
0.025
0.000
1.0
0.8
0.6
0.4
0.2
0.0
ICM
Random
Count
RND
Disagreement
Low-pass OU actions
High-pass OU actions