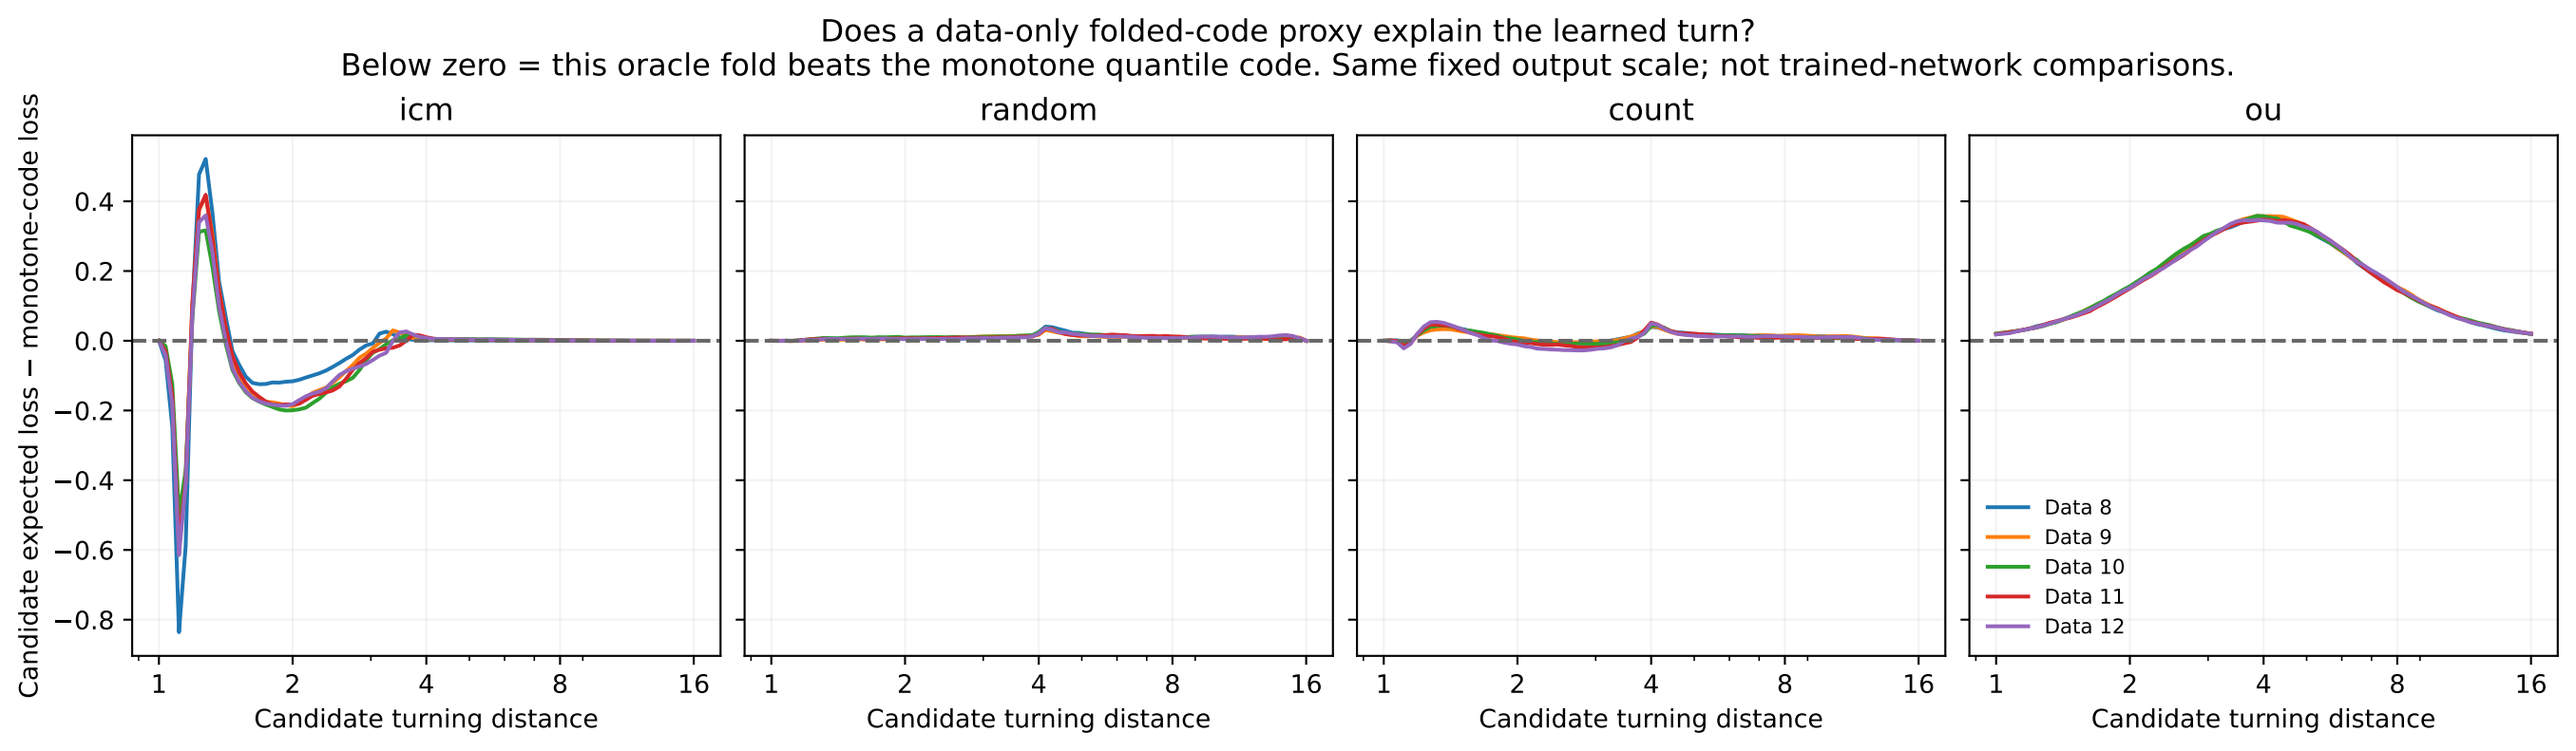
Does a data-only folded-code proxy explain the learned turn?
Below zero = this oracle fold beats the monotone quantile code. Same fixed output scale; not trained-network comparisons.
icm
random
count
ou
Candidate expected loss − monotone-code loss
0.4
0.2
0.0
−0.2
−0.4
−0.6
−0.8
1
2
4
8
16
Candidate turning distance
Data 8
Data 9
Data 10
Data 11
Data 12