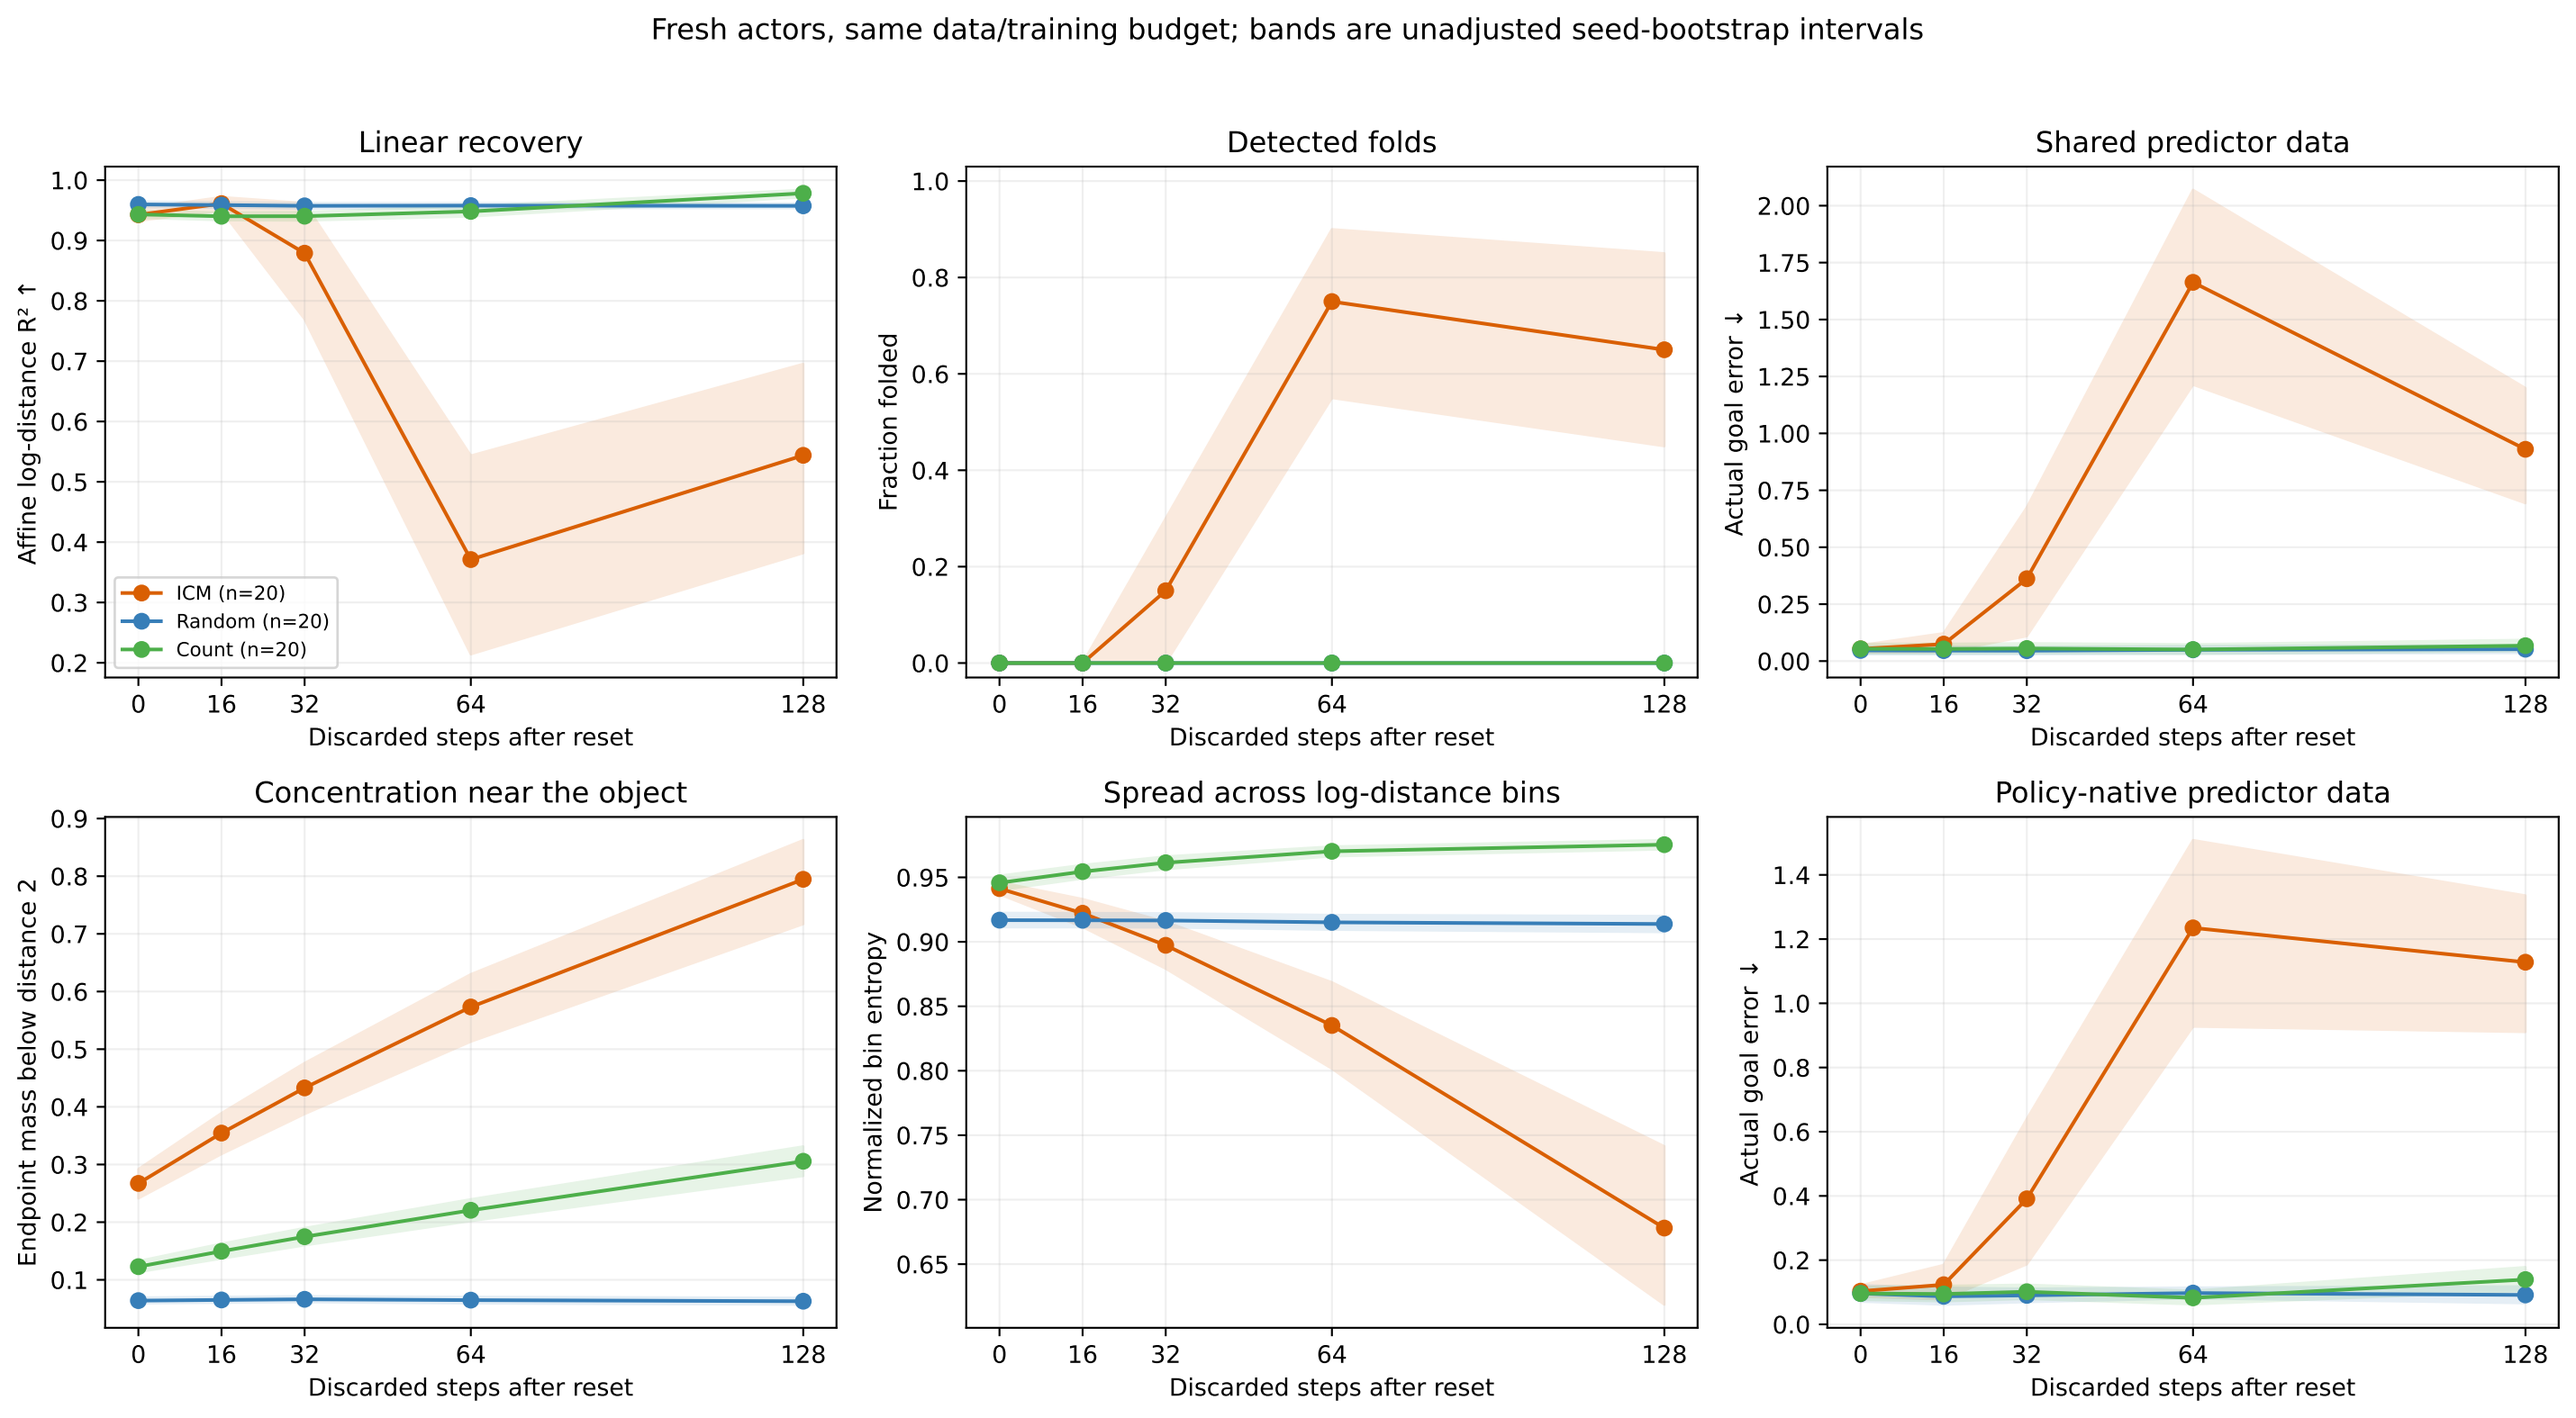
Fresh actors, same data/training budget; bands are unadjusted seed-bootstrap intervals
Linear recovery
Affine log-distance R² ↑
ICM (n=20)
Random (n=20)
Count (n=20)
Discarded steps after reset
Detected folds
Fraction folded
Discarded steps after reset
Shared predictor data
Actual goal error ↓
Discarded steps after reset
Concentration near the object
Endpoint mass below distance 2
Discarded steps after reset
Spread across log-distance bins
Normalized bin entropy
Discarded steps after reset
Policy-native predictor data
Actual goal error ↓
Discarded steps after reset