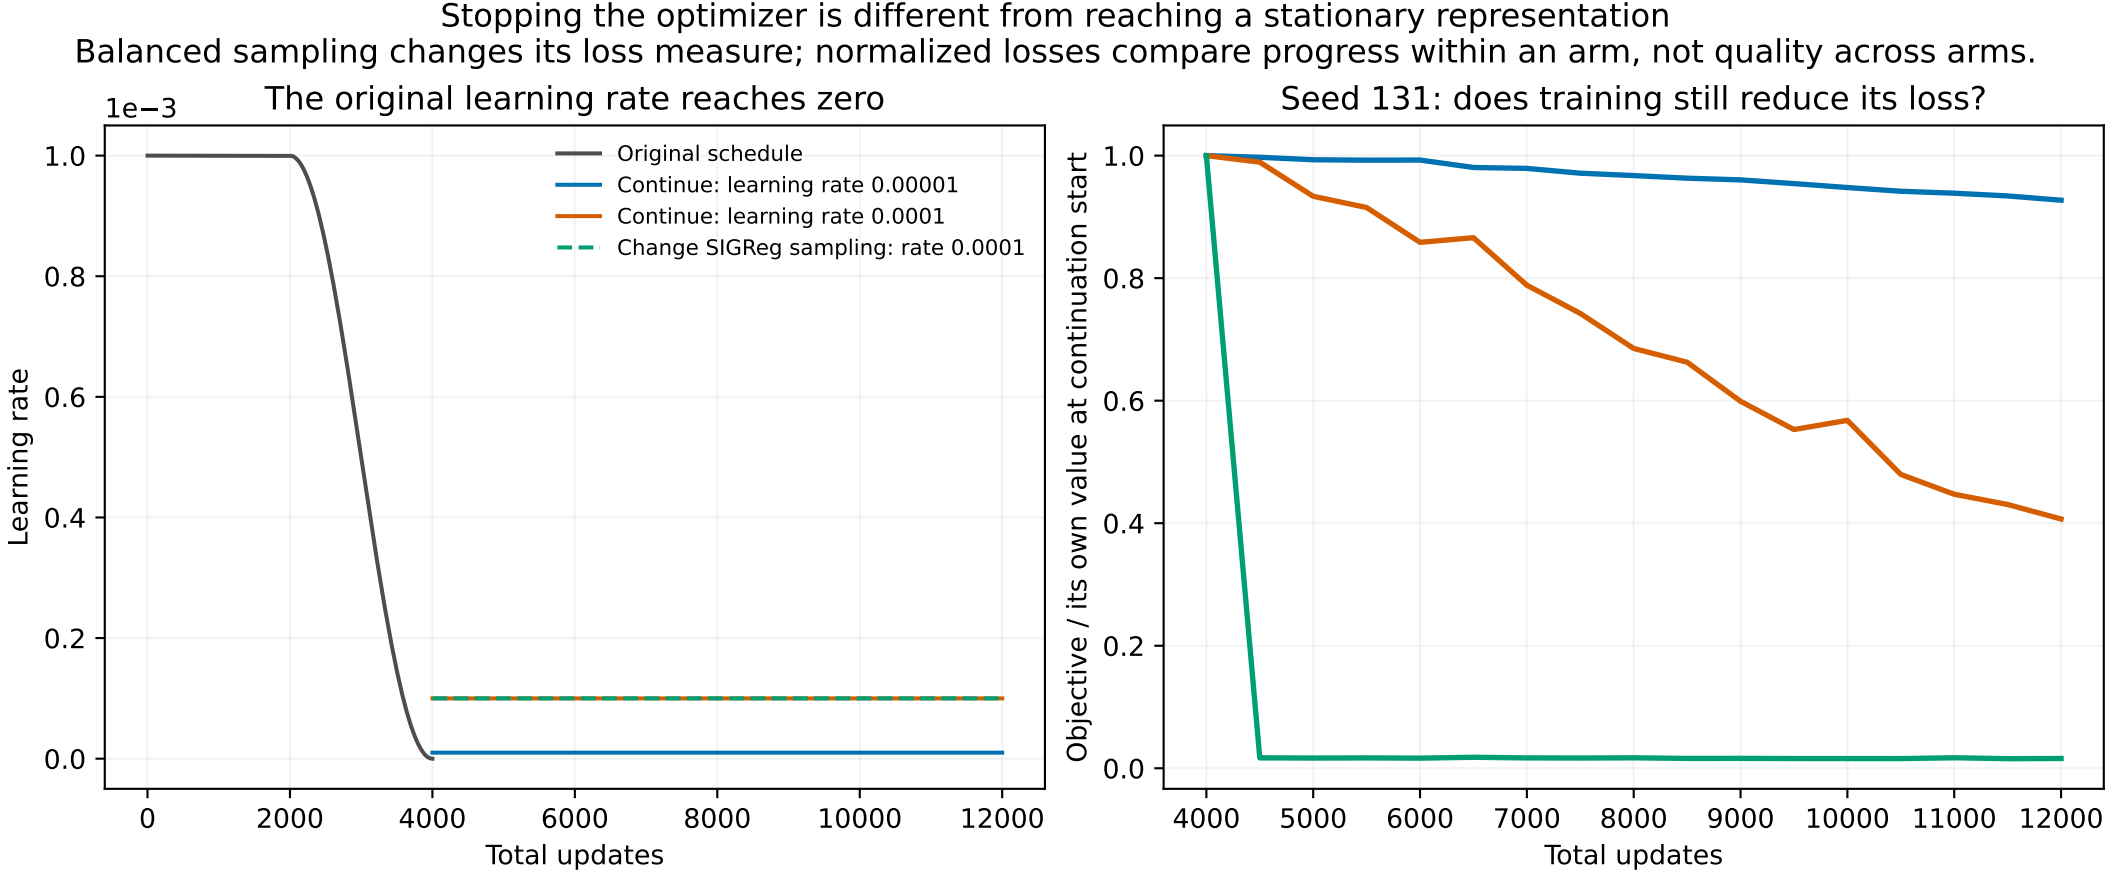
Stopping the optimizer is different from reaching a stationary representation
Balanced sampling changes its loss measure; normalized losses compare progress within an arm, not quality across arms.
The original learning rate reaches zero
1e−3
Learning rate
Total updates
Original schedule
Continue: learning rate 0.00001
Continue: learning rate 0.0001
Change SIGReg sampling: rate 0.0001
Seed 131: does training still reduce its loss?
Objective / its own value at continuation start
Total updates